

# Language Modeling

## Unit-II

**Syed Rameem Zahra**  
(Assistant Professor)  
Department of CSE, NSUT

# What is a Language Model in NLP?

- A language model learns to predict the probability of a sequence of words.
- It's a statistical tool that analyzes the pattern of human language for the prediction of words, by estimating the relative likelihood of different phrases.
- The models are prepared for the prediction of words by learning the features and characteristics of a language.
- Language models are used in speech recognition, machine translation, part-of-speech tagging, parsing, Optical Character Recognition, handwriting recognition, information retrieval, summarization, spell correction, and many other daily tasks.

# Challenges with Language Modeling

- Formal languages (like a programming language) are precisely defined, but Natural language isn't designed, it evolves according to the convenience and learning of an individual.
- There are several terms in natural language that can be used in a number of ways, which introduces ambiguity but can still be understood by humans.

# Some Common Examples of Language Models

- **Speech Recognition**

- Voice assistants such as Siri and Alexa are examples of how language models help machines in processing speech audio.

- **Machine Translation**

- Google Translator and Microsoft Translate are examples of how NLP models can help in translating one language to another.

- **Sentiment Analysis**

- This helps in analyzing the sentiments behind a phrase. This use case of NLP models is used in products that allow businesses to understand a customer's intent behind opinions or attitudes expressed in the text.

- **Text Suggestions**

- Google services such as Gmail or Google Docs use language models to help users get text suggestions while they compose an email or create long text documents, respectively.

- **Parsing Tools**

- Parsing involves analyzing sentences or words that comply with syntax or grammar rules. Spell checking tools are perfect examples of language modelling and parsing.

# Types of Language Models

- **Statistical Language Models:**

- These models use traditional statistical techniques like **N-grams**, Hidden Markov Models (**HMM**) and certain linguistic rules to learn the probability distribution of words.

- **Neural Language Models:**

- These are new players in the NLP town and have surpassed the statistical language models in their effectiveness. They use different kinds of Neural Networks to model language.

# Goal of Probabilistic Language Modelling

- To calculate the probability of a sentence or sequence of words:

- $P(W) = P(w_1, w_2, \dots, w_n)$

- This is the **Joint Probability**

- It can be calculated using **Conditional probability**:

- $P(w_5 | w_1, w_2, w_3, w_4)$

$$P(w_5 | w_1, w_2, w_3, w_4) = \frac{\text{count}(w_1, w_2, w_3, w_4, w_5)}{\text{count}(w_1, w_2, w_3, w_4)}$$

- E.g. for two words: X, Y, we have:

- $P(X|Y) = P(X,Y)/P(Y)$

- then,  $P(X,Y) = P(X|Y) P(Y)$

- Similarly, for three words:  $P(X,Y,Z) = P(X) P(Y|X) P(Z|X,Y)$

- This method is used when we have short/limited sentence.

- For longer sentences, we can use **Markov assumption**:

- $P(w_n | w_1, w_2, w_3, \dots, w_{n-1}) \sim P(w_n | w_{n-1})$  or  $P(w_n | w_{n-2}, w_{n-1})$

# N-Gram Model

- This is one of the simplest approaches to language modelling.
- Here, a probability distribution for a sequence of 'n' is created, where 'n' can be any number and defines the size of the gram (or sequence of words being assigned a probability).
- If  $n=4$ , a gram may look like: "can you help me".
- Basically, 'n' is the amount of context that the model is trained to consider.
- There are different types of N-Gram models such as **unigrams**, **bigrams**, **trigrams**, etc.
- The intuition of the n-gram model is that **instead of computing the probability of a word given its entire history, we can approximate the history by just the last few words.**

## **(Uni-) 1-gram model**

- The simplest case of Markov assumption is case when the size of prefix is one.
- This will provide us with grammar that only consider one word. As a result it produces a set of unrelated words.
- It actually would generate sentences with random word order.



# Bigram Model

- Approximates the probability of a word given all the previous words by using only the conditional probability of the preceding word.
- we consider a 2-word (**tandem**) bigrams correlations
- In other words, instead of computing the probability
  - $P(\text{the}|\text{Walden Pond's water is so transparent that})$
- we approximate it with the probability
  - $P(\text{the}|\text{that})$

# How to estimate these bigram or n-gram probabilities?

- An intuitive way to estimate probabilities is called **maximum likelihood estimation** or **MLE**.
- We get Maximum likelihood estimate for the parameters of an n-gram model by getting counts from a corpus, and normalizing the counts so that they lie between 0 and 1.
- ***For example***, to compute a particular bigram probability of a word  $\mathbf{w}_n$  given a previous word  $\mathbf{w}_{n-1}$ , we'll compute the count of the bigram  $\mathbf{C}(\mathbf{w}_{n-1}\mathbf{w}_n)$  and normalize by the sum of all the bigrams that share the same first word  $\mathbf{w}_{n-1}$

**Note:** A **corpus** is a collection of authentic text or audio organized into datasets.

## Example-1

- We have a mini-corpus of three sentences:
- $\langle s \rangle$  I am Sam  $\langle /s \rangle$
- $\langle s \rangle$  Sam I am  $\langle /s \rangle$
- $\langle s \rangle$  I do not like green eggs and ham  $\langle /s \rangle$
- Here are the calculations for some of the bigram probabilities from this corpus:

$$\begin{aligned} P(I | \langle s \rangle) &= \frac{2}{3} = .67 & P(\text{Sam} | \langle s \rangle) &= \frac{1}{3} = .33 & P(\text{am} | I) &= \frac{2}{3} = .67 \\ P(\langle /s \rangle | \text{Sam}) &= \frac{1}{2} = 0.5 & P(\text{Sam} | \text{am}) &= \frac{1}{2} = .5 & P(\text{do} | I) &= \frac{1}{3} = .33 \end{aligned}$$

- In practice it's more common to use **trigram** models, which condition on the previous two words rather than the previous word, or **4-gram** or even **5-gram** models, when there is sufficient training data.

## Example-2: Bi-gram probabilities

- What is the most probable next word predicted by the model for the following word sequence:

### Given Corpus

<S> I am Henry </S>

<S> I like college </S>

<S> Do Henry like college </S>

<S> Henry I am </S>

<S> Do I like Henry </S>

<S> Do I like college </S>

<S> I do like Henry </S>

Word	Frequency
<S>	7
</S>	7
I	6
am	2
Henry	5
like	5
college	3
do	4

## Example-2: Bi-gram probabilities (Contd...)

1) <S> Do ?

<S> I am Henry </S>  
 <S> I like college </S>  
 <S> Do Henry like college </S>  
 <S> Henry I am </S>  
 <S> Do I like Henry </S>  
 <S> Do I like college </S>  
 <S> I do like Henry </S>

Word	Frequency
<S>	7
</S>	7
I	6
am	2
Henry	5
like	5
college	3
do	4

Next word prediction probability  $W_{i-1} = \text{do}$

Next word	Probability $\frac{\text{count}(w_{i-1}, w_i)}{\text{count}(w_{i-1})}$
$P(</S>   \text{do})$	0/4
$P(<I>   \text{do})$	2/4
$P(<\text{am}>   \text{do})$	0/4
$P(<\text{Henry}>   \text{do})$	1/4
$P(<\text{like}>   \text{do})$	1/4
$P(<\text{college}>   \text{do})$	0/4
$P(\text{do}   \text{do})$	0/4

**I is more probable**

## Example-2: Bi-gram probabilities (Contd...)

2) <S> I like Henry ?

<S> I am Henry </S>  
<S> I like college </S>  
<S> Do Henry like college </S>  
<S> Henry I am </S>  
<S> Do I like Henry </S>  
<S> Do I like college </S>  
<S> I do like Henry </S>

Word	Frequency
<S>	7
</S>	7
I	6
am	2
Henry	5
like	5
college	3
do	4

Next word prediction probability  $W_{i-1} = \text{Henry}$

Next word	Probability Next Word = $\frac{N}{D} = \frac{\text{count}(w_{i-1}, w_i)}{\text{count}(w_{i-1})}$
$P(</S>   \text{Henry})$	3/5
$P(<I>   \text{Henry})$	1/5
$P(<\text{am}>   \text{Henry})$	0
$P(<\text{Henry}>   \text{Henry})$	0
$P(<\text{like}   \text{Henry})$	1/5
$P(<\text{college}   \text{Henry})$	0
$P(<\text{do}   \text{Henry})$	0

</S> is more probable

## Example-2: Bi-gram probabilities (Contd...)

Which of the following sentence is better. i.e. Gets a higher probability with this model.  
Use Bi-gram

<S> I am Henry </S>  
<S> I like college </S>  
<S> Do Henry like college </S>  
<S> Henry I am </S>  
<S> Do I like Henry </S>  
<S> Do I like college </S>  
<S> I do like Henry </S>

Word	Frequency
<S>	7
</S>	7
I	6
am	2
Henry	5
like	5
college	3
do	4

1. <S> I like college </S>

<S> like college </S>=?

$$\begin{aligned} &= P(I \mid <S>) \times P(\text{like} \mid I) \times P(\text{college} \mid \text{like}) \times P(</S> \mid \text{college}) \\ &= 3/7 \times 3/6 \times 3/5 \times 3/3 = 9/70 = 0.13 \end{aligned}$$


2. <S> Do I like Henry </S>

$$\begin{aligned} &= P(\text{do} \mid <S>) \times P(I \mid \text{do}) \times P(\text{like} \mid I) \times P(\text{Henry} \mid \text{like}) \times P(</S> \mid \text{Henry}) \\ &= 3/7 \times 2/4 \times 3/6 \times 2/5 \times 3/5 = 9/350 = 0.0257 \end{aligned}$$

**ANS: First statement is more probable**

# Publicly available corpora

- **Gutenberg Project** providing with text format of some books.
- **Google** also released a publicly available corpus, trillion word corpus with over 13 million unique words.





# N-gram Language Model

- **Advantages:**

- Easy to understand and implement
- Conversion from one type of gram to another is easy.

- **Disadvantages:**

- Underflow due to multiplication of probabilities
  - Solution: Use Log (which will add probabilities)
- Zero probability problem
  - Solution: Use Laplace smoothing

# Using Log to solve underflow problem

Which of the following sentence is better. i.e. Gets a higher probability with Bi-gram model.

<S> I am Henry </S>

<S> I like college </S>

<S> Do Henry like college </S>

<S> Henry I am </S>

<S> Do I like Henry </S>

<S> Do I like college </S>

<S> I do like Henry </S>

Word	Frequency
<S>	7
</S>	7
I	6
am	2
Henry	5
like	5
college	3
do	4

**First statement is more probable**

## 1. <S> I like college </S>

$$=P(I | <S>) \times P(\text{like} | I) \times P(\text{college} | \text{like}) \times P(</S> | \text{college})$$

$$=3/7 \times 3/6 \times 3/5 \times 3/3 = 9/70 = \mathbf{0.13}$$

$$= \log(3/7) + \log(3/6) + \log(3/5) + \log(3/3) = \mathbf{-2.0513}$$

## 2. <S> Do I like Henry </S>

$$=P(\text{do} | <S>) \times P(I | \text{do}) \times P(\text{like} | I) \times P(\text{Henry} | \text{like}) \times P(</S> | \text{Henry})$$

$$=3/7 \times 2/4 \times 3/6 \times 2/5 \times 3/5 = 9/350 = \mathbf{0.0257}$$

$$= \log(3/7) + \log(2/4) + \log(3/6) + \log(2/5) + \log(3/5) = \mathbf{-3.6607}$$

# Zero Probability Problem

<S> I am Henry </S>

<S> I like college </S>

<S> Do Henry like college </S>

<S> Henry I am </S>

<S> Do I like Henry </S>

<S> Do I like college </S>

<S> I do like Henry </S>

Word	Frequency
<S>	7
</S>	7
I	6
am	2
Henry	5
like	5
college	3
do	4

**Second statement is more probable**

## 1. <S> like college </S>

$$=P(\text{like} \mid \text{<S>}) \times P(\text{college} \mid \text{like}) \times P(\text{</S>} \mid \text{college})$$

$$=0/7 \times 3/5 \times 3/3 = \mathbf{0}$$



## 2. <S> Do I like Henry </S>

$$=P(\text{do} \mid \text{<S>}) \times P(\text{I} \mid \text{do}) \times P(\text{like} \mid \text{I}) \times P(\text{Henry} \mid \text{like}) \times P(\text{</S>} \mid \text{Henry})$$

$$=3/7 \times 2/4 \times 3/6 \times 2/5 \times 3/5 = 9/350 = \mathbf{0.0257}$$

# Smoothing

- To keep a language model from assigning zero probability to these unseen events, we'll have to shave off a bit of probability mass from some more frequent events and give it to the events we've never seen.
- This modification is called smoothing (or discounting).
- There are many ways to do smoothing, and some of them are:
  - Add-1 smoothing (Laplace Smoothing)
  - Add-k smoothing,
  - Backoff
  - Kneser-Ney smoothing.

Bigram	Frequency	Bigram	Frequency
CS 421	8	CS 421	7
CS 590	5	CS 590	5
CS 594	2	CS 594	2
CS 521	0 🤔	CS 521	1 🧐

# Laplace Smoothing

- Add one to all n-gram counts before they are normalized into probabilities.
- Not the highest-performing technique for language modeling, but useful method for text classification

$$P(w_i) = \frac{c_i}{N} \rightarrow P_{\text{Laplace}}(w_i) = \frac{c_i+1}{N+V}$$

# Laplace Smoothing

<S> I am Henry </S>  
<S> I like college </S>  
<S> Do Henry like college </S>  
<S> Henry I am </S>  
<S> Do I like Henry </S>  
<S> Do I like college </S>  
<S> I do like Henry </S>

Word	Frequency
<S>	7
</S>	7
I	6
am	2
Henry	5
like	5
college	3
do	4

Unique words are : <S>, </S>, I, Henry do, like, am, college

**Total unique words: 8**

But we exclude <S> as it never comes in bi-gram calculations

**Total unique words: 7**



**Give the following bi-gram probabilities estimated by Laplace model.**

# Applying Laplace Smoothing

1. <S> like college </S>

$$=P(\text{like} \mid \text{<S>}) \times P(\text{college} \mid \text{like}) \times P(\text{</S>} \mid \text{college})$$

$$=(0+1)/(7+7) \times (3+1)/(5+7) \times (3+1)/(3+7)$$

$$=1/14 \times 4/12 \times 4/10$$

$$=\mathbf{0.0095}$$

2. <S> Do I like Henry </S>

$$=P(\text{do} \mid \text{<S>}) \times P(\text{I} \mid \text{do}) \times P(\text{like} \mid \text{I}) \times P(\text{Henry} \mid \text{like}) \times P(\text{</S>} \mid \text{Henry})$$

$$=(3+1)/(7+7) \times (2+1)/(4+7) \times (3+1)/(6+7) \times (2+1)/(5+7) \times (3+1)/(5+7)$$

$$=4/14 \times 3/11 \times 4/13 \times 3/12 \times 4/12$$

$$=\mathbf{0.0020}$$

# Add-K Smoothing

- Rather than adding one to each count, add a fractional count (0.5, 0.05, 0.01 etc.)
- The value of K can be optimized on a validation set

$$P(w_i) = \frac{c_i}{N} \rightarrow P_{\text{Add-K}}(w_i) = \frac{c_i + k}{N + kV}$$
$$P(w_n | w_{n-1}) = \frac{c(w_{n-1}w_n)}{c(w_{n-1})} \rightarrow P_{\text{Add-K}}(w_n | w_{n-1}) = \frac{c(w_{n-1}w_n) + k}{c(w_{n-1}) + kV}$$



# Backoff and Interpolation

- Add-K smoothing is useful for some tasks, but still tends to be suboptimal for language modeling.
- Other techniques are:
  - **Backoff:** we use the trigram if the evidence is sufficient, otherwise we use the bigram, otherwise the unigram.
    - In other words, we only “back off” to a lower-order n-gram if we have zero evidence for a higher-order n-gram.
  - **Interpolation:** we always mix the probability estimates from all the n-gram estimators, weighing and combining the trigram, bigram, and unigram counts.

# Simple Linear Interpolation

- we combine different order n-grams by linearly interpolating all the models.

$$\begin{aligned}\hat{P}(w_n|w_{n-1}w_{n-2}) &= \lambda_1 P(w_n|w_{n-1}w_{n-2}) \\ &\quad + \lambda_2 P(w_n|w_{n-1}) \\ &\quad + \lambda_3 P(w_n)\end{aligned}\qquad \sum_i \lambda_i = 1$$

## **Language Model Evaluation:**

Language model is better if it is assigning a high probability to the real, frequently observed and grammatical sentence over false, rarely observed and ungrammatical sentences.

Two different criteria for evaluation

### **1) Extrinsic    2) Intrinsic**

#### **Extrinsic Evaluation**

It evaluate the language model when solving a specific task.

For e.g. Speech recognition accuracy, Machine translation accuracy, Spelling correction accuracy Compare 2 (or more) models, and check which works best.

#### **Disadvantage:**

- Expensive
- Time consuming

## **Intrinsic Evaluation:**

The language model is best when it predicts an unseen test set.

### **Definition of Perplexity:**

It is the inverse probability of the test data which is normalized by the number of words.

Lower the value of perplexity: **Better Model**

More value of perplexity: **Confused for prediction**

# Perplexity

The language model is best when it predicts an unseen test set.

## Definition of Perplexity:

It is the inverse probability of the test data which is normalized by the number of words.

$$PP(w) = P(w_1, w_2, w_3, \dots, w_N)^{-\frac{1}{N}}$$

$$PP(w) = \left( \prod_i \frac{1}{P(w_i | w_1, w_2, \dots, w_{i-1})} \right)^{\frac{1}{N}} \quad PP(w) = \left( \prod_i \frac{1}{P(w_i | w_{i-1})} \right)^{\frac{1}{N}}$$

Lower the value of perplexity: **Better Model**

More value of perplexity: **Confused for prediction**

# Example

**WSJ Corpus**

**Training:** 38 million words

**Test:** 1.5 million words

N-gram Order	Unigram	Bigram	Trigram
Perplexity	962	170	109

**Perplexity for Bigram <S> I like college </S>**

$$=P(I \mid <S>) \times P(\text{like} \mid I) \times P(\text{college} \mid \text{like}) \times P(</S> \mid \text{college})$$

$$=3/7 \times 3/6 \times 3/5 \times 3/3 = 9/70 = \mathbf{0.13}$$

$$\mathbf{PP(w) = (1/0.13)^{1/4} = 1.67}$$

**Perplexity for Trigram <S> I like college </S>**

$$P(w) = P(\text{like} \mid <S> I) \times P(\text{college} \mid I \text{ like}) \times P(</S> \mid \text{like college})$$

$$P(w) = 1/3 \times 2/3 \times 3/3 = 2/9 = \mathbf{0.22}$$

$$\mathbf{PP(w) = (1/0.22)^{1/3} = 1.66}$$

# Text Classification

- Text Classification (Text Categorization) is the task of assigning a label or categorization category to an entire text or document.
- Some of common text categorization tasks are:
  - **Sentiment analysis**
    - extraction of sentiment, positive or negative orientation that writer expresses toward an object.
  - **Spam detection**
    - binary classification task of assigning an email to one of the two classes spam or not-spam.
  - **Authorship identification**
    - determining a text's author.
  - **Age/gender identification**
    - determining a text's author characteristics like gender and age.
  - **Language Identification**
    - finding the language of a text.

# Text Classification: Definition

- Text classification can be defined as follows:
- Input:
  - a document  $d$
  - a fixed set of classes  $C = \{c_1, c_2, \dots, c_n\}$
- Output:
  - a predicted class  $c$  belongs-to  $C$



# Classification Methods: Hand-Coded Rules

- The goal of classification is to take a single observation, extract some useful features, and thereby classify the observation into one of a set of discrete classes.
- One method for classifying text is to use hand-written rules.
- Rules based on combinations of words or other features
  - spam: black-list-address OR (“dollars” AND “have been selected”)
- Accuracy can be high if rules carefully are refined by experts.
- But building and maintaining these hand-written rules can be expensive.
  - Rules can be fragile.
  - It may require domain knowledge.

# Classification Methods: Supervised Machine Learning

- Most cases of text classification in language processing are done via supervised machine learning methods.
- The goal of a supervised machine learning algorithm is to learn how to map from a new observation to a correct output.
- Input:
  - a document  $d$
  - a fixed set of classes  $C = \{c_1, c_2, \dots, c_n\}$
  - a training set of  $m$  hand-labeled documents  $(d_1, c_1), \dots, (d_m, c_m)$
- Output:
  - a learned classifier model:  $d \rightarrow c$

# Classification Methods: Supervised Machine Learning

- Our goal is to learn a classifier that is capable of mapping from a new document  $d$  to its correct class  $c$  belongs-to  $C$ .
- A probabilistic classifier additionally will tell us the probability of the observation being in the class.
- Generative classifiers like Naive Bayes build a model of how a class could generate some input data.
  - Given an observation, they return the class most likely to have generated the observation.
- Discriminative classifiers like logistic regression instead learn what features from the input are most useful to discriminate between the different possible classes.
- Some classifiers:
  - Naïve Bayes
  - Logistic regression
  - Support-vector machines
  - k-Nearest Neighbors, ...

# Text Classification: Evaluation

- In order to evaluate **how good is our classifier**, we can use different evaluation metrics.
- In evaluation, we compare the **test set results** of our classifier with **gold labels** (the human labels for the test set documents).
- As a result of this comparison, first we build a **contingency table (or confusion matrix)** before calculate our evaluation metrics:

		gold standard labels		
		gold positive	gold negative	
system output labels	system positive	true positive	false positive	precision = $\frac{tp}{tp+fp}$
	system negative	false negative	true negative	
		recall = $\frac{tp}{tp+fn}$		accuracy = $\frac{tp+tn}{tp+fp+tn+fn}$

# Text Classification: Evaluation

- **Accuracy** is percentage of all the observations our system labeled correctly.
- **Precision** measures percentage of items that system detected that are in fact positive.
- **Recall** measures percentage of items actually present in the input that were correctly identified by the system.
- Precision and Recall, unlike Accuracy, emphasize true positives.
  - Looking only one of them can be misleading.
  - $tp=1$   $fp=0$   $fn=99$ , then, Precision = 100% (while Recall=1%)
  - $tp=1$   $fp=99$   $fn=0$ , then, Recall= 100% (while Precision=1%)
- **F-measure** is a single metric that incorporates aspects of both precision and recall.

# Text Classification: Evaluation

$$F_{\beta} = \frac{(\beta^2 + 1) * P * R}{\beta^2 * P + R}$$

The  $\beta$  parameter differentially weights the importance of recall and precision.

- values of  $\beta > 1$  favor recall, while values of  $\beta < 1$  favor precision.

The most frequently used metric, and is called  $F_{\beta=1}$  or just  $F_1$ :

$$F_1 = \frac{2 * P * R}{P + R}$$

# Text Classification with more than two classes

		<i>gold labels</i>			
		urgent	normal	spam	
<i>system output</i>	urgent	8	10	1	$\text{precision}_u = \frac{8}{8+10+1}$
	normal	5	60	50	$\text{precision}_n = \frac{60}{5+60+50}$
	spam	3	30	200	$\text{precision}_s = \frac{200}{3+30+200}$
		$\text{recall}_u = \frac{8}{8+5+3}$	$\text{recall}_n = \frac{60}{10+60+30}$	$\text{recall}_s = \frac{200}{1+50+200}$	$\text{Accuracy} = \frac{\sum_i c_{ii}}{\sum_j \sum_i c_{ij}}$

# Microaveraging and Macroaveraging

- In order to derive a single metric that tells us how well the system is doing, we can combine these values in two ways.
  - In **macroaveraging**, compute performance for each class, and then average over classes.
  - In **microaveraging**, collect decisions for all classes into a single contingency table, and then compute precision and recall from that table.

Class 1: Urgent

	true urgent	true not
system urgent	8	11
system not	8	340

precision =  $\frac{8}{8+11} = .42$

Class 2: Normal

	true normal	true not
system normal	60	55
system not	40	212

precision =  $\frac{60}{60+55} = .52$

Class 3: Spam

	true spam	true not
system spam	200	33
system not	51	83

precision =  $\frac{200}{200+33} = .86$

Pooled

	true yes	true no
system yes	268	99
system no	99	635

microaverage  
precision =  $\frac{268}{268+99} = .73$

macroaverage  
precision =  $\frac{.42+.52+.86}{3} = .60$




# Cross-validation

- we randomly choose a training and test set division of our data, train our classifier, and then compute the error rate on the test set.
- Then we repeat with a different randomly selected training set and test set.
- We do sampling process 10 times and average these 10 runs to get an average error rate.
- This is called 10-fold cross-validation.

# Part-of-Speech Tagging

- Given a text of english, identify the parts of speech of each word



# Part-of-Speech Tagging

- **Open class words** (content words):
  - Nouns, verbs, adjectives, adverbs.
  - They refer to objects, actions and features in the world.
  - They are open class because new words are added all the time.
- **Closed class words:**
  - Pronouns, determiners, prepositions, connectives,...
  - They are limited.
  - Mostly functional: to tie the concepts of a sentence together.

# Part-of-Speech Tagging

N	noun	chair, bandwidth, pacing
V	verb	study, debate, munch
ADJ	adj	purple, tall, ridiculous
ADV	adverb	unfortunately, slowly,
P	preposition	of, by, to
PRO	pronoun	I, me, mine
DET	determiner	the, a, that, those

# POS Tagging: Choosing a target

- For POS tagging, we need to choose a standard set.
- E.g., We could choose a very coarse targets like N, V, Adj, Adv etc.
- A commonly used set is: “UPenn TreeBank” tagset, which contains 45 tags.

# UPenn TreeBank POS tagset

Tag	Description	Example	Tag	Description	Example
CC	Coordin. Conjunction	<i>and, but, or</i>	SYM	Symbol	<i>+, %, &amp;</i>
CD	Cardinal number	<i>one, two, three</i>	TO	"to"	<i>to</i>
DT	Determiner	<i>a, the</i>	UH	Interjection	<i>ah, oops</i>
EX	Existential 'there'	<i>there</i>	VB	Verb, base form	<i>eat</i>
FW	Foreign word	<i>mea culpa</i>	VBD	Verb, past tense	<i>ate</i>
IN	Preposition/sub-conj	<i>of, in, by</i>	VBG	Verb, gerund	<i>eating</i>
JJ	Adjective	<i>yellow</i>	VCN	Verb, past participle	<i>eaten</i>
JJR	Adj., comparative	<i>bigger</i>	VBP	Verb, non-3sg pres	<i>eat</i>
JJS	Adj., superlative	<i>wildest</i>	VBZ	Verb, 3sg pres	<i>eats</i>
LS	List item marker	<i>1, 2, One</i>	WDT	Wh-determiner	<i>which, that</i>
MD	Modal	<i>can, should</i>	WP	Wh-pronoun	<i>what, who</i>
NN	Noun, sing. or mass	<i>llama</i>	WPS	Possessive wh-	<i>whose</i>
NNS	Noun, plural	<i>llamas</i>	WRB	Wh-adverb	<i>how, where</i>
NNP	Proper noun, singular	<i>IBM</i>	\$	Dollar sign	<i>\$</i>
NNPS	Proper noun, plural	<i>Carolinas</i>	#	Pound sign	<i>#</i>
PDT	Predeterminer	<i>all, both</i>	"	Left quote	<i>(' or ")</i>
POS	Possessive ending	<i>'s</i>	"	Right quote	<i>(' or ")</i>
PRP	Personal pronoun	<i>I, you, he</i>	(	Left parenthesis	<i>( [, (, {, &lt;)</i>
PRP\$	Possessive pronoun	<i>your, one's</i>	)	Right parenthesis	<i>( ], ), }, &gt;)</i>
RB	Adverb	<i>quickly, never</i>	,	Comma	<i>,</i>
RBR	Adverb, comparative	<i>faster</i>	.	Sentence-final punc	<i>(. ! ?)</i>
RBS	Adverb, superlative	<i>fastest</i>	:	Mid-sentence punc	<i>(: ; ... - -)</i>
RP	Particle	<i>up, off</i>			

# POS Tagging is hard???

- A word in a sentence may have multiple POS tags depending on the context.
  - E.g., For the word “Back”, we have:
  - The back door: *back/JJ* -> Adjective
  - On my back: *back/NN* -> Noun
  - Win the voters back: *back/RB* -> Adverb
- It has been seen that mostly we have 2 to 3 tags for many words (Ambiguity problem)
- We can use any valid corpus to find the highest probability of a word for tagging it to a particular POS tag, which designing a model.
  - E.g.: Some words may only be nouns like arrow
  - Some words are ambiguous like flies
  - Probability may help us if one tag is more likely than another.
  - Also the local context can be used.

# POS Tagging Approaches

- **Rule based approach:**

- Assign each word in the input sentence a list of potential POS tags.
- Then, scale down the list to a single tag using hand-written rules.

- **Statistical tagging:**

- Get a training corpus of tagged text, learn the transformation rules from most frequently tags (e.g. TBL (Transformation Based Learning) Tagger).
- Find the most likely sequence of tags for a sequence of words using probability.




# Generative and Conditional Models

- **Generative (Joint) Model:**


- Generate the observed data from hidden stuff, i.e., put a probability over the observation given the class:  $P(d,c)$  in terms of  $P(d|c)$
- E.g., Naive bayes classification, Hidden Markov Models etc.

- **Conditional (Discriminative) Models:**

- Take the data as given and put a probability over hidden structure given the data:  $P(c|d)$
- E.g., Logistic regression, max. Entropy models, SVMs, Perceptron etc.




Naive Bayes



Logistic Regression

# Generative Model: Hidden Markov Model

- A hidden Markov model (HMM) is a statistical Markov model in which the system being modeled is assumed to be a Markov process — call it  $\mathbf{X}$  with unobservable ("hidden") states.
- As part of the definition, HMM requires that there be an observable process  $\mathbf{Y}$  whose outcomes are "influenced" by the outcomes of  $\mathbf{X}$  in a known way.
- Since  $\mathbf{X}$  cannot be observed directly, the goal is to learn about  $\mathbf{X}$  by observing  $\mathbf{Y}$ .
- HMM has an additional requirement that the outcome of  $\mathbf{Y}$  at time  $t = t_0$  must be "influenced" exclusively by the outcome of  $\mathbf{X}$  at time  $t = t_0$  and that the outcomes of  $\mathbf{X}$  and  $\mathbf{Y}$  at  $t < t_0$  must not affect the outcome of  $\mathbf{Y}$  at  $t = t_0$ .
- Similar to N-Gram models
- Model the text as a sequence
- For ngrams, we modeled the probability of each word conditioned on the previous  $n-1$  words.
- Here, we model each tag conditioned on previous tags
- Still uses **Markov assumption**: only look back a few tags



$X$  — states  
 $y$  — possible observations  
 $a$  — state transition probabilities  
 $b$  — output probabilities

# Generative Model: Hidden Markov Model

- We want the most likely tag sequence for a sequence of words:

$$p(\langle S \rangle t_1 t_2 \dots t_n \langle E \rangle \mid \langle S \rangle w_1 w_2 \dots w_n \langle E \rangle)$$

Remember that order matters!

- For simplicity, we'll write this as  $t_1^n = \langle t_1, t_2, \dots, t_n \rangle$
- So we want

$$\hat{t}_1^n = \operatorname{argmax}_{t_1^n} p(\langle S \rangle t_1 t_2 \dots t_n \langle E \rangle \mid \langle S \rangle w_1 w_2 \dots w_n \langle E \rangle)$$

but it's hard to estimate

- Like for ngrams, use Bayes rule:

$$\begin{aligned}\hat{t}_1^n &= \operatorname{argmax}_{t_1^n} \frac{p(\langle S \rangle w_1 w_2 \dots w_n \langle E \rangle \mid \langle S \rangle t_1 t_2 \dots t_n \langle E \rangle) \cdot p(\langle S \rangle t_1 t_2 \dots t_n \langle E \rangle)}{p(\langle S \rangle w_1 w_2 \dots w_n \langle E \rangle)} \\ &= \operatorname{argmax}_{t_1^n} p(\langle S \rangle w_1 w_2 \dots w_n \langle E \rangle \mid \langle S \rangle t_1 t_2 \dots t_n \langle E \rangle) \cdot p(\langle S \rangle t_1 t_2 \dots t_n \langle E \rangle)\end{aligned}$$

- Two major independence assumptions:

– Like ngrams, assume probability of a sequence is dependent only on recent past:

$$\begin{aligned}p(\langle S \rangle t_1 t_2 \dots t_n \langle E \rangle) &\approx p(t_1 \mid \langle S \rangle) \cdot p(t_2 \mid t_1) \cdot p(t_2 \mid t_2) \cdot \dots \cdot p(t_n \mid t_{n-1}) \cdot p(\langle E \rangle \mid t_n) \\ &= \prod_{i=1}^n p(t_i \mid t_{i-1})\end{aligned}$$

# Generative Model: Hidden Markov Model

- Also assume word is only dependent on its tag:

$$\begin{aligned} p(\langle S \rangle w_1 w_2 \dots w_n \langle E \rangle \mid \langle S \rangle t_1 t_2 \dots t_n \langle E \rangle) &\approx p(w_1 \mid t_1) \cdot p(w_2 \mid t_2) \cdot \dots \cdot p(w_n \mid t_n) \\ &= \prod_{i=1}^n p(w_i \mid t_i) \end{aligned}$$

- Together:

$$\begin{aligned} \hat{t}_1^n &= \operatorname{argmax}_{t_1^n} p(\langle S \rangle t_1 t_2 \dots t_n \langle E \rangle \mid \langle S \rangle w_1 w_2 \dots w_n \langle E \rangle) \\ &= \operatorname{argmax}_{t_1^n} p(\langle S \rangle w_1 w_2 \dots w_n \langle E \rangle \mid \langle S \rangle t_1 t_2 \dots t_n \langle E \rangle) \cdot p(\langle S \rangle t_1 t_2 \dots t_n \langle E \rangle) \\ &\approx \operatorname{argmax}_{t_1^n} \prod_{i=1}^n p(w_i \mid t_i) \cdot \prod_{i=1}^n p(t_i \mid t_{i-1}) \\ &= \operatorname{argmax}_{t_1^n} \prod_{i=1}^n p(w_i \mid t_i) \cdot p(t_i \mid t_{i-1}) \end{aligned}$$

# Estimating Parameters: MLE in Hidden Markov Model

Two probability distributions to estimate:

- Transitions: probability of a tag, given previous tag,  $p(t_i | t_{i-1})$
- Emissions: probability of a word, given its tag,  $p(w_i | t_i)$

MLE

- MLE estimation is just like before (naïve Bayes, ngrams, ...): normalized counts
- Transitions:  $p(t_i | t_{i-1}) = \frac{C(t_{i-1} \ t_i)}{\sum_x C(t_{i-1} \ x)} = \frac{C(t_{i-1} \ t_i)}{C(t_{i-1})}$
- Emissions:  $p(w_i | t_i) = \frac{C(t_i, w_i)}{\sum_x C(t_i, x)} = \frac{C(t_i, w_i)}{C(t_i)}$

# Example

Example dataset (punctuation excluded for simplicity):

```
<S>|<S> the|D man|N walks|V the|D dog|N <E>|<E>
<S>|<S> the|D dog|N runs|V <E>|<E>
<S>|<S> the|D dog|N walks|V <E>|<E>
<S>|<S> the|D man|N walks|V <E>|<E>
<S>|<S> a|D man|N saw|V the|D dog|N <E>|<E>
<S>|<S> the|D cat|N walks|V <E>|<E>
```

Some probabilities:

$$\bullet p(t_i = D \mid t_{i-1} = N) = \frac{C(N, D)}{\sum_x C(N, x)} = \frac{0}{8} = 0.0$$

$$\bullet p(t_i = V \mid t_{i-1} = N) = \frac{C(N, V)}{\sum_x C(N, x)} = \frac{6}{8} = 0.75 \quad \bullet p(w_i = \text{dog} \mid t_i = N) = \frac{C(N, \text{dog})}{\sum_x C(N, x)} = \frac{4}{8} = 0.50$$

$$\bullet p(w_i = \text{the} \mid t_i = N) = \frac{C(N, \text{the})}{\sum_x C(N, x)} = \frac{0}{8} = 0.0$$

*Note: We can add smoothing techniques also (as discussed earlier)*

# Three Tasks for HMM

- **Likelihood:** Given a tagged sequence, determine its likelihood
- **Decoding:** Given an untagged sequence, determine the best tag sequence for it
- **Learning:** Given an untagged sequence, and a set of tags, learn the HMM parameters



# Likelihood of a tagged sentence

We can compute the likelihood of a particular sequence of tags for a sentence:


- $p(t_1 \dots t_n \mid w_1 \dots w_n) = \prod_{i=1}^n p(w_i \mid t_i) \cdot p(t_i \mid t_{i-1})$

Example: "the|D dog|N walks|V"

$$\begin{aligned} p(t_1 \dots t_n \mid w_1 \dots w_n) &\approx \prod_{i=1}^n p(w_i \mid t_i) \cdot p(t_i \mid t_{i-1}) \\ &= p(D \mid \langle S \rangle) \cdot p(the \mid D) \cdot p(N \mid D) \cdot p(dog \mid N) \cdot p(V \mid N) \cdot p(walks \mid V) \cdot p(\langle E \rangle \mid V) \end{aligned}$$



# Example: Disambiguation of “race”



- Using HMM and available corpus, we can find most likely sequence of tags using probability values.

# Issues with Markov Model Tagging

- Missing probabilities for **unknown words**.
  - ***Solution:*** Use morphological cues (Capitalization or suffixes etc.) to assign a more calculated guess.
- **Limited context** may not be sufficient for correct tagging.
  - ***Solution:*** use higher order HMM (like 2nd order or 3rd order etc.) and combine various N-gram models.

# Maximum Entropy Model: Discriminative Model

- Uses a combination of heterogeneous set of features to create a probabilistic model, which is able to select a correct PoS tag for a current word, e.g.:
  - Whether the next word is **to**.
  - Whether one of the last 5 words is a **preposition**. etc.

$$p_{\lambda}(y|x) = \frac{1}{Z_{\lambda}(x)} \exp\left(\sum_i \lambda_i f_i(x, y)\right)$$

where

- $Z_{\lambda}(x)$  is a normalizing constant given by

$$Z_{\lambda}(x) = \sum_y \exp\left(\sum_i \lambda_i f_i(x, y)\right)$$

- $\lambda_i$  is a weight given to a feature  $f_i$
- $x$  denotes an observed datum and  $y$  denotes a class

- **Principles of Max. Entropy Model:** Given a collection of facts (features), choose a model which is consistent (Uniform) with all the facts.

# Features in Max. Entropy Model

- Features encode elements of the context  $x$  for predicting tag  $y$
- Context  $x$  is taken around the word  $w$ , for which a tag  $y$  is to be predicted
- Features are binary valued functions, e.g.,

$$f(x, y) = \begin{cases} 1 & \text{if } isCapitalized(w) \& y = NNP \\ 0 & \text{otherwise} \end{cases}$$

## Examples of features:

### *Example: Named Entities*

- LOCATION (in Arcadia)
- LOCATION (in Québec)
- DRUG (taking Zantac)
- PERSON (saw Sue)

### *Example Features*

- $f_1(x, y) = [y = LOCATION \wedge w_{-1} = "in" \wedge isCapitalized(w)]$
- $f_2(x, y) = [y = LOCATION \wedge hasAccentedLatinChar(w)]$
- $f_3(x, y) = [y = DRUG \wedge ends(w, "c")]$

# PoS Tagging in Max. Entropy Model

- $W = w_1 \dots w_n$  - words in the corpus (observed)
- $T = t_1 \dots t_n$  - the corresponding tags (unknown)

Tag sequence candidate  $\{t_1, \dots, t_n\}$  has conditional probability:

$$P(t_1, \dots, t_n | w_1 \dots, w_n) = \prod_{i=1}^n p(t_i | x_i)$$

- The context  $x_i$  also includes previously assigned tags for a fixed history.

# Entropy in Max Entropy Model

- It measures the uncertainty (surprise) of a distribution.
- For an event  $x$  with probability of occurrence  $p_x$ , Entropy =  $\log (1/p_x)$
- Entropy  $H$  for a random variable  $X$  with probability distribution  $P$  is given as:

$$E_p \left[ \log_2 \frac{1}{p_x} \right] = - \sum_x p_x \log_2 p_x$$

- So, in Max. Entropy Model, we choose a model with max. Entropy, subjected to feature-based constraints.
- We will start from a uniform distribution (because it has max. Entropy) and the add constraints, which will decrease the entropy and make it closer to the given data.

# Example

Consider the maximum entropy model for POS tagging, where you want to estimate  $P(\text{tag}|\text{word})$ . In a hypothetical setting, assume that *tag* can take the values *D*, *N* and *V* (short forms for Determiner, Noun and Verb). The variable *word* could be any member of a set *V* of possible words, where *V* contains the words *a*, *man*, *sleeps*, as well as additional words. The distribution should give the following probabilities

- $P(D|a) = 0.9$
- $P(N|man) = 0.9$
- $P(V|sleeps) = 0.9$
- $P(D|\text{word}) = 0.6$  for any word other than *a*, *man* or *sleeps*
- $P(N|\text{word}) = 0.3$  for any word other than *a*, *man* or *sleeps*
- $P(V|\text{word}) = 0.1$  for any word other than *a*, *man* or *sleeps*

It is assumed that all other probabilities, not defined above could take any values such that  $\sum_{\text{tag}} P(\text{tag}|\text{word}) = 1$  is satisfied for any word in *V*.

- Define the features of your maximum entropy model that can model this distribution. Mark your features as  $f_1, f_2$  and so on. Each feature should have the same format as explained in the class. [Hint: 6 Features should make the analysis easier]
- For each feature  $f_i$ , assume a weight  $\lambda_i$ . Now, write expression for the following probabilities in terms of your model parameters
  - ▶  $P(D|cat)$
  - ▶  $P(N|laughs)$
  - ▶  $P(D|man)$
- What value do the parameters in your model take to give the distribution as described above. (i.e.  $P(D|a) = 0.9$  and so on. You may leave the final answer in terms of equations)

## Example Contd...

- **F1** =  $F1(x,y) = [\text{word} = \text{'a'} \ \& \ \text{tag} = \text{'D'}]$
- **F2** =  $F2(x,y) = [\text{word} = \text{'man'} \ \& \ \text{tag} = \text{'N'}]$
- **F3** =  $F3(x,y) = [\text{word} = \text{'sleeps'} \ \& \ \text{tag} = \text{'V'}]$
- **F4** =  $F4(x,y) = [\text{word} \in V' \ \& \ \text{tag} = \text{'D'}]$ , Where  $V' = V - \{\text{a,man,sleeps}\}$ ;  
V is Vocabulary
- **F5** =  $F5(x,y) = [\text{word} \in V' \ \& \ \text{tag} = \text{'N'}]$
- **F6** =  $F6(x,y) = [\text{word} \in V' \ \& \ \text{tag} = \text{'V'}]$
- Now,  **$P(D|cat) = e^{(\sum \lambda_i F_i)} / Z$** 
  - $\sum \lambda_i F_i = \lambda_1 * 0 + \lambda_2 * 0 + \lambda_3 * 0 + \lambda_4 * 1 + \lambda_5 * 0 + \lambda_6 * 0 = \lambda_4$
  - To calculate Z, we need to calculate  $P(N|cat)$  and  $P(V|cat)$
  - $P(N|cat) = e^{\lambda_5} / Z$  and  $P(V|cat) = e^{\lambda_6} / Z$ ; Also,  
 $P(D|cat) + P(N|cat) + P(V|cat) = 1 \Rightarrow Z = e^{\lambda_4} + e^{\lambda_5} + e^{\lambda_6}$
  - Hence,  **$P(D|cat) = e^{\lambda_4} / (e^{\lambda_4} + e^{\lambda_5} + e^{\lambda_6})$**
- Similarly, we can calculate  **$P(N|laugh)$**  and  **$P(D|man)$**



## Example Contd...

- $P(D|a) = e^{\lambda_1}/Z$ , to calculate  $Z$  here, we have:
  - $P(V|a)=e^0/Z=1/Z$     $P(N|a)=e^0/Z=1/Z$
  - Hence,  $P(D|a) = e^{\lambda_1}/(e^{\lambda_1}+2) = 0.9$
- Similarly, we can calculate equation for other given constraints and get values of  $\lambda$ s, which will represent the overall Max. entropy model for the given problem.